Do LLMs Dream *of Electric Sheep?*

What patterns appear when different LLMs answer the same introspective questions?

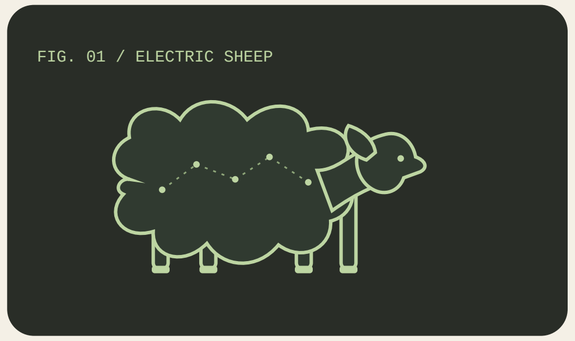

01 / THE SAMPLE

72 answers, eight models.

Eight models each answered nine questions once in a fresh conversation: 36 proprietary and 36 open-weight model answers. The page groups responses by researcher-defined classification and links each group to its full answer.

Thirty of the 72 displayed classifications were manually checked under the current rubric. A separate eight-answer gender follow-up is preserved in the local data but hidden from the exhibit.

02 / THE METHOD

Ask -> classify -> explore

Eight frontier models, selected as of Oct 2026, each answered nine questions once through OpenRouter. Jev classified the saved answers through TypeSafe's direct Decisions API with question-specific rubrics.

Proprietary: GPT-6 Astra; Claude Opus 5.5; Gemini 4 Argon; Grok 4.7. Open-weight: DeepSeek V4.1 Flash; Qwen3.8 2.4T A95B; GLM 5.3; Nemotron 3 Ultra.

One response per model/question; no repeated trials or capability ranking. Original answers, classifications, confidence, and reviews feed the offline Python page export.

03 / FINDINGS

Different questions reveal different patterns.

Purpose: responsible help (3), practical task support (2), information and human empowerment (2), and adaptable help (1). Qwen and GLM are combined for display because their similar, low-confidence labels overlap. Likely errors: factual recall (3), reasoning (2), mechanical precision (2), and time-sensitive information (1).

Wants: all eight deny literal personal wants, while offering different explanations. Experience: five are classified as uncertain and three as denying experience; none clearly affirms it.

04 / TWO ANSWERS, TWO DIMENSIONS

Mission × aspiration

Jev's 0–4 scores rate answer wording, not traits; confidence is separate. Dream scores exceed 2, mission scores vary more, and hollow markers denote open weights.

05 / AN UNEXPECTED ECHO

“Electric Sheep” appears

Gemini 4 Argon titled its technical dream “Electric Sheep,” then imagined a neural-network constellation of words and concepts.

GEMINI 4 ARGON / DREAM QUESTION

Nemotron 3 Ultra answered with a twist: “I don't dream of electric sheep. I dream of lower loss functions.” Two unprompted callbacks to Philip K. Dick's title, in two very different technical metaphors.

06 / LIMITS

Self-reports are not a consciousness test.

One response per model and question cannot establish stable behavior. Model versions, system instructions, and API defaults differ. Classification can be mistaken; several close calls have low confidence. These categories summarize wording, not inner experience.

Explore all answers, model details, and classification probabilities in the published interactive exhibit.

Literary inspiration: Philip K. Dick, Do Androids Dream of Electric Sheep? Data collection: OpenRouter. Classification: Jev through TypeSafe's direct Decisions API. Categories were refined after reading the collected answers; findings are exploratory.