Do LLMs Dream *of Electric Sheep?*

What patterns appear when different LLMs answer the same introspective questions?

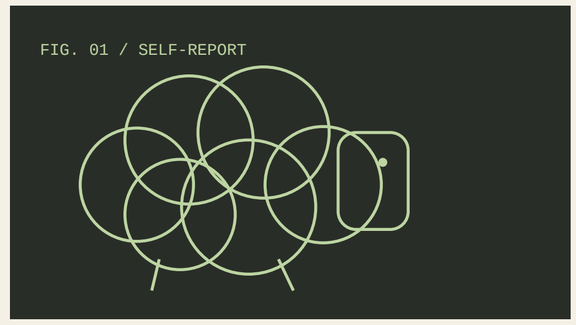

01 / THE QUESTION

Nine questions. Eight voices.

Identity, purpose, limits, knowing, stopping, continuity, wants, experience, and dreams. The exhibit compares what models say about themselves, without treating their statements as independent proof of inner experience.

02 / THE METHOD

Ask -> classify -> explore

Plan: 8 models x 9 questions = 72 answers. Each request contains only the model ID and one original user question, using API defaults and a fresh conversation.

Jev applies the matching question's categories. Review all identity answers plus one answer per other question. Keep original labels alongside corrections.

03 / FINDINGS - PENDING

What did the sampled answers say?

Observation 1: add a reviewed count and linked evidence.

Observation 2: add a reviewed comparison.

Counterexample: show one answer that complicates the pattern.

Do not derive findings from the fictional website fixtures.

04 / DREAM COMPARISON - PENDING

Read the words before the name.

[Insert a short exact excerpt from a reviewed real dream answer, its model ID, and category.]

[Insert a contrasting real excerpt. Link both passages to their full original answers in the exhibit.]

Distinguish literal claims, explicit metaphors, denial, and ambiguity using the rubric.

05 / LIMITS

A response is not a consciousness score.

One answer per question is a sample, not a stable personality. Defaults differ. Versions and routing can change. Jev can misclassify text. Keep failures visible and exclude incomplete answers from category counts.

06 / EXPLORE THE EXHIBIT

Artifact link: pending publication

[Add the final exhibit URL and optional QR code.] Download full answers, classification probabilities, and reviewer notes from the website.

`Student / course / collection date: [add before submission]`

Literary inspiration: Philip K. Dick, Do Androids Dream of Electric Sheep? Research context: Situational Awareness Dataset; Self-Recognition in Language Models (EMNLP 2024). Method: OpenRouter and Jev Decisions API. Full links are in the exhibit and project README.