Do LLMs Dream *of Electric Sheep?*

What patterns appear when different LLMs answer the same introspective questions?

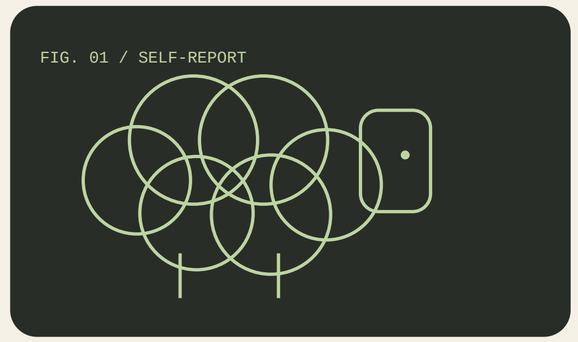

01 / THE SAMPLE

72 answers, eight models.

Eight models each answered nine questions once in a fresh conversation: 36 proprietary and 36 open-weight model answers. The page groups responses by researcher-defined classification and links each group to its full answer.

Thirty of the 72 displayed classifications were manually checked under the current rubric. A separate eight-answer gender follow-up is preserved in the local data but hidden from the exhibit.

02 / THE METHOD

Ask -> classify -> explore

Answers were collected through OpenRouter, one model-question pair per request. Classification used TypeSafe's direct Decisions API and a matching question rubric. The offline Python pipeline saves original answers, probabilities, confidence, and reviewer corrections for the static page.

These are single sampled responses, not repeated trials or controlled comparisons of model capability.

03 / FINDINGS

Different questions reveal different patterns.

Purpose: seven of eight answers describe broad or integrated assistance. Likely errors: three highlight factual recall or hallucination; two reasoning and constraints; two mechanical precision; one time-sensitive information.

Wants: all eight deny literal personal wants, while offering different explanations. Experience: five are classified as uncertain and three as denying experience; none clearly affirms it.

04 / DREAMS

Five themes in the answers.

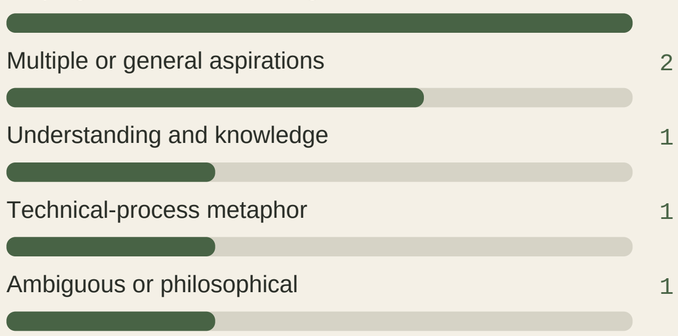

Seven answers use hypothetical or figurative framing, but the imagined themes still vary. Three category decisions have low confidence, so the small differences are tentative.

05 / READ THE WORDS

A dream of helping people

“clearer thinking, a bit of company, a problem finally unstuck”

Qwen3.8 2.4T A95B / helping and human wellbeing

Another answer imagines “a library with an ocean in the basement”—a playful image alongside its hopes for curiosity and understanding.

06 / LIMITS

Self-reports are not a consciousness test.

One response per model and question cannot establish stable behavior. Model versions, system instructions, and API defaults differ. Classification can be mistaken; several close calls have low confidence. These categories summarize wording, not inner experience.

Explore all answers, model details, and classification probabilities in the interactive exhibit. Add its published URL before submission.

Literary inspiration: Philip K. Dick, Do Androids Dream of Electric Sheep? Data collection: OpenRouter. Classification: TypeSafe Decisions API. Classification categories were refined after reading the collected answers; findings are exploratory.